

LINGUISTIC AND COMPUTATIONAL SEMANTICS*

Brian Cantwell Smith
XEROX Palo Alto Research Center
3333 Coyote Hill Road, Palo Alto, CA 94304

ABSTRACT

We argue that because the very concept of computation rests on notions of interpretation, the semantics of natural languages and the semantics of computational formalisms are in the deepest sense the same subject. The attempt to use computational formalisms in aid of an explanation of natural language semantics, therefore, is an enterprise that must be undertaken with particular care. We describe a framework for semantical analysis that we have used in the computational realm, and suggest that it may serve to underwrite computationally-oriented linguistic semantics as well. The major feature of this framework is the explicit recognition of both the declarative and the procedural import of meaningful expressions; we argue that whereas these two viewpoints have traditionally been taken as alternative, any comprehensive semantical theory must account for how both aspects of an expression contribute to its overall significance.

We have argued elsewhere¹ that the distinguishing mark of those objects and processes we call *computational* has to do with attributed semantics: we humans find computational processes coherent exactly because we attach semantical significance to their behaviour, ingredients, and so forth. Put another way, computers, on our view, are those devices that we understand by *deploying our linguistic faculties*. For example, the reason that a calculator is a computer, but a car is not, is that we take the ingredients of the calculator to be *symbolic* (standing, in this particular case, for numbers and functions and so forth), and understand the interactions and organisation of the calculator in terms of that interpretation (this part *divides*, this part *represents the sum*, and so on). Even though by and large we are able to produce an explanation of the behaviour that does not rest on external semantic attribution (this is the *formality condition* mentioned by Fodor, Haugeland, and others²), we nonetheless speak, when we use computational terms, in terms of this semantics. These semantical concepts rest at the foundations of the discipline: the particular organisations that computers have — their computational *raison d'être* — emerge not only from their mechanical structure but also from their semantic interpretability. Similarly, the terms of art employed in computer science — *program*, *compiler*, *implementation*, *interpreter*, and so forth — will ultimately be definable only with reference to this attributed semantics; they will not, on our view, ever be found reducible to non-semantical predicates.³


This is a ramifying and problematic position, which we cannot defend here.⁴ We may simply note, however, the overwhelming evidence in favour of a semantical approach manifested by everyday computational language. Even the simple view of computer science as the study of *symbol manipulation*⁵ reveals this bias. Equally telling is the fact that programming languages are called *languages*. In addition, language-derived concepts like *name* and *reference* and *semantics* permeate computational jargon (to say nothing of *interpreter*, *value*, *variable*, *memory*, *expression*, *identifier* and so on) — a fact that would be hard to explain if semantics were not crucially involved. It is not just that in discussing computation we *use* language; rather, in discussing computation we use words that suggest that we are also talking *about* linguistic phenomena.

The question we will focus on in this paper, very briefly, is this: if computational artefacts are fundamentally linguistic, and if, therefore, it is appropriate to analyse them in terms of formal theories of semantics (it is apparent that this is a widely held view), then what is the proper relationship between the so-called *computational* semantics that results, and more standard *linguistic* semantics (the discipline that studies people and their natural languages: how we mean, and what we are talking about, and all of that good stuff)? And furthermore, what is it to *use* computational models to *explain* natural language semantics, if the computational models are themselves in need of semantical analysis? On the face of it, there would seem to be a certain complexity that should be sorted out.

In answering these questions we will argue approximately as follows: in the limit computational semantics and linguistic semantics will coincide, at least in underlying conception, if not in surface detail (for example some issues, like ambiguity, may arise in one case and not in the other). Unfortunately, however, as presently used in computer science the term "semantics" is given such an operational cast that it distracts attention from the human attribution of significance to computational structures.⁶ In contrast, the most successful models of natural language semantics, embodied for example in standard model theories and even in Montague's program, have concentrated almost exclusively on *referential* or *denotational* aspects of declarative sentences. Judging only by surface use, in other words, computational semantics and linguistic semantics appear almost orthogonal in *concern*, even though they are of course similar in *style* (for example they both use meta-theoretic mathematical techniques — functional composition, and so forth — to recursively specify the semantics of complex expressions from a


given set of primitive atoms and formation rules). It is striking, however, to observe two facts. First, computational semantics is being pushed (by people and by need) more and more towards declarative or referential issues. Second, natural language semantics, particularly in computationally-based studies, is focusing more and more on pragmatic questions of use and psychological import. Since computational linguistics operates under the computational hypothesis of mind, psychological issues are assumed to be modelled by a field of computational structures and the state of a processor running over them; thus these linguistic concerns with "use" connect naturally with the "operational" flavour of standard programming language semantics. It seems not implausible, therefore — we betray our caution with the double negative — that a unifying framework might be developed.

It will be the intent of this paper to present a specific, if preliminary, proposal for such a framework. First, however, some introductory comments. In a general sense of the term, *semantics* can be taken as the study of the relationship between entities or phenomena in a *syntactic domain* S and corresponding entities in a *semantic domain* D , as pictured in the following diagram.



We call the function mapping elements from the first domain into elements of the second an *interpretation function* (to be sharply distinguished⁷ from what in computer science is called an *interpreter*, which is a different beast altogether). Note that the question of whether an element is syntactic or semantic is a function of the point of view; the syntactic domain for one interpretation function can readily be the semantic domain of another (and a semantic domain may of course include its own syntactic domain).

Not all relationships, of course, count as semantical; the "grandmother" relationship fits into the picture just sketched, but stakes no claim on being semantical. Though it has often been discussed what constraints on such a relationship characterise genuinely semantical ones (compositionality or recursive specifiability, and a certain kind of formal character to the syntactic domain, are among those typically mentioned), we will not pursue such questions here. Rather, we will complicate our diagram as follows, so as to enable us to characterise a rather large class of computational and linguistic formalisms:



N_1 and N_2 are intended to be *notational* or *communicational expressions*, in some externally observable and consensually established medium of interaction, such as strings of characters, streams of words, or sequences of display images on a computer terminal. The relationship θ is an interpretation function mapping notations into *internal elements* of some process over which the primary semantical and processing regimens are defined. In first-

order logic, S_1 and S_2 would be something like abstract derivation tree types of first-order formulae; if the diagram were applied to the human mind, under the hypothesis of a formally encoded mentalese, S_1 and S_2 would be tokens of internal mentalese, and θ would be the function computed by the "linguistic" faculty (on a view such as that of Fodor⁸). In adopting these terms we mean to be speaking *very* generally; thus we mean to avoid, for example, any claim that *tokens* of English are *internalised* (a term we will use for θ) into recognisable tokens of mentalese. In particular, the proper account of θ for humans could well simply describe how the field of mentalese structures, in some configuration, is transformed into some other configuration, upon being presented with a particular English sentence; this would still count, on our view, as a theory of θ .

In contrast, ϕ is the interpretation function that makes explicit the standard denotational significance of linguistic terms, relating, we may presume, expressions in S to the world of discourse. The relationship between my mental token for T. S. Eliot, for example, and the poet himself, would be formulated as part of ϕ . Again, we speak very broadly; ϕ is intended to manifest what, paradigmatically, expressions are *about*, however that might best be formulated (ϕ includes for example the interpretation functions of standard model theories). ψ , in contrast, relates some internal structures or states to others — one can imagine it specifically as the formally computed derivability relationship in a logic, as the function computed by the primitive language processor in a computational machine (i.e., as LISP's EVAL), or more generally as the function that relates one configuration of a field of symbols to another, in terms of the modifications engendered by some internal processor computing over those states. (ϕ and ψ are named, for mnemonic convenience, by analogy with *philosophy* and *psychology*, since a study of ϕ is a study of the relationship between expressions and the world — since philosophy takes you "out of your mind", so to speak — whereas a study of ψ is a study of the internal relationships between symbols, all of which, in contrast, are "within the head" of the person or machine.)

Some simple comments. First, N_1 , N_2 , S_1 , S_2 , D_1 , and D_2 need not all necessarily be distinct: in a case where S_1 is a self-referential designator, for example, D_1 would be the same as S_1 ; similarly, in a case where ψ computed a function that was *designation-preserving*, then D_1 and D_2 would be identical. Secondly, we need not take a stand on which of ψ and ϕ has a prior claim to being the *semantics* of S_1 . In standard logic, ψ (i.e., derivability: \vdash) is a relationship, but is far from a function, and there is little tendency to think of it as *semantical*; a study of ψ is called *proof* theory. In computational systems, on the other hand, ψ is typically much more constrained, and is also, by and large, analysed mathematically in terms of functions and so forth, in a manner much more like standard model theories. Although in this author's view it seems a little far-fetched to call the internal relationships (the "use" of a symbol) *semantical*, it is nonetheless true that we are interested in characterising both, and it is unnecessary to express a preference. For discussion, we will refer to the ψ -semantics of a symbol or expression as its *declarative import*, and refer to its ϕ -semantics as its *procedural consequence*. We have heard it said in other quarters that "procedural" and "declarative" theories of semantics are contenders;⁹ to the extent that we have been able to make sense of these notions, it appears that we need both.

It is possible to use this diagram to characterise a variety of standard formal systems. In the standard models of the λ -calculus, for example, the designation function Φ takes λ -expressions onto functions; the procedural regimen Ψ , usually consisting of α - and β -reductions, can be shown to be Φ -preserving. Similarly, if in a standard predicate logic we take Φ to be (the inverse of the) satisfaction relationship, with each element of S being a sentence or set of sentences, and elements of D being those possible worlds in which those sentences are true, and similarly take Ψ as the derivability relationship, then soundness and completeness can be expressed as the equation $\Psi(S_1, S_2) = [D_1 \subseteq D_2]$. As for all formal systems (these presumably subsume the computational ones), it is crucial that Ψ be specifiable independent of Φ . The λ -calculus and predicate logic systems, furthermore, have no notion of a processor with state; thus the appropriate Ψ involves what we may call local procedural consequence, relating a simple symbol or set of symbols to another set. In a more complex computational circumstance, as we will see below, it is appropriate to characterise a more complex full procedural consequence involving not only simple expressions, but fuller encodings of the state of various aspects of the computational machine (for example, at least environments and continuations in the typical computational case¹⁰).

An important consequence of the analysis illustrated in the last figure is that it enables one to ask a question not typically asked in computer science, about the (Φ) semantic character of the function computed by Ψ . Note that questions about soundness and completeness in logic are exactly questions of this type. In separate research,¹¹ we have shown, by subjecting it to this kind of analysis, that computational formalisms can be usefully analysed in these terms as well. In particular, we demonstrated that the universally accepted LISP evaluation protocol is semantically confused, in the following sense: sometimes it preserves Φ (i.e. $\Phi(\Psi(S)) = \Phi(S)$), and sometimes it embodies Φ (i.e., $\Psi(S) = \Phi(S)$). The traditional LISP notion of evaluation, in other words, conflates *simplification* and *reference* relationships, to its peril (in that report we propose some LISP dialects in which these two are kept strictly separate). The current moral, however, is merely that our approach allows the question of the semantical import of Ψ to be asked.

As well as considering LISP, we may use our diagram to characterise the various linguistically oriented projects carried on under the banner of "semantics". Model theories and formal theories of language (we include Tarski and Montague in one sweep) have concentrated primarily on Φ . Natural language semantics in some quarters¹² focuses on Θ — on the translation into an internal medium — although the question of what aspects of a given sentence must be preserved in such a translation are of course of concern (no translator could ignore the salient properties, semantical and otherwise, of the target language, be it mentalese or predicate logic, since the endeavour would otherwise be without constraint). Lewis (for one) has argued that the project of articulating Θ — an endeavour he calls *markerese* semantics — cannot really be called semantics at all,¹³ since it is essentially a translation relationship, although it is worth noting that Θ in computational formalisms is not always trivial, and a case can at least be made that many superficial aspects of natural language use, such as the resolution of indexicals, may be resolved at this stage (if for example you say *I am warm* then I may internalise your use of the first person pronoun into my internal name for you).

Those artificial intelligence researchers working in knowledge representation, perhaps without too much distortion, can be divided into two groups: a) those whose primary semantical allegiance is to Φ , and who (perhaps as a consequence) typically use an encoding of first-order logic as their representation language, and b) those who concern themselves primarily with Ψ , and who therefore (legitimately enough) reject logic as even suggestive (Ψ in logic — derivability — is a relatively unconstrained relationship, for one thing; secondly, the relationship between the entailment relationship, to which derivability is a hopeful approximation, and the proper " Ψ " of rational belief revision, is at least a matter of debate¹⁴).

Programming language semantics, for reasons that can at least be explored, if not wholly explained, have focused primarily on Ψ , although in ways that tend to confuse it with Φ . Except for PROLOG, which borrows its Φ straight from a subset of first-order logic, and the LISPs mentioned earlier,¹⁵ we have never seen a semantical account of a programming language that gave independent accounts of Φ and Ψ . There are complexities, furthermore, in knowing just what the proper treatment of general languages should be. In a separate paper¹⁶ we argue that the notion *program* is inherently defined as a set of expressions whose (Φ) semantic domain includes *data structures* (and set-theoretic entities built up over them). In other words, in a computational process that deals with finance, say, the *general* data structures will likely designate individuals and money and relationships among them, but the terms in that part of the process called a *program* will not designate these people and their money, but will instead designate: *the data structures that designate people and money* (plus of course relationships and functions over those data structures). Even on a *declarative* view like ours, in other words, the appropriate semantic domain for programs is built up over data structures — a situation strikingly like the standard semantical accounts that take abstract records or locations or whatever as elements of the otherwise mathematical domain for programming language semantics. It may be that this fact that all base terms in programs are *meta-syntactic* that has spawned the confusion between operations and reference in the computational setting.

Although the details of a general story remain to be worked out, the LISP case mentioned earlier is instructive, by way of suggestion as to how a more complete computational theory of language semantics might go. In particular, because of the context relativity and non-local effects that can emerge from processing a LISP expression, Φ is not specifiable in a strict compositional way. Ψ — when taken to include the broadest possible notion that maps entire configurations of the field of symbols and of the processor itself onto other configurations and states — is of course recursively specifiable (the same fact, in essence, as saying that LISP is a deterministic formal calculus). A pure characterisation of Ψ without a concomitant account of Φ , however, is unmotivated — as empty as a specification of a derivability relationship would be for a calculus for which no semantics had been given. Of more interest is the ability to specify what we call a general significance function Σ , that recursively specifies Ψ and Φ together (this is what we were able to do for LISP). In particular, given any expression s_1 , any configuration of the rest of the symbols, and any state of the processor, the function Σ will specify the configuration and state that would result (i.e., it will specify the use of s_1), and also the relationship to the world that the whole signifies. For example,

given a LISP expression of the form $(+ 1 (\text{PROG} (\text{SETQ} A 2) A))$, Σ would specify that the whole expression designated the number three, that it would return the numeral "3", and that the machine would be left in a state in which the binding of the variable A was changed to the numeral "2". A modest result; what is important is merely a) that *both* declarative import and procedural significance must be reconstructed in order to tell a full story about LISP; and b) that they must be formulated together.

Rather than pursue this view in detail, it is helpful to set out several points that emerge from analyses developed within this framework:

- a. In most programming languages, Θ can be specified compositionally and independently of Φ or Ψ — this amounts to a formal statement of Fodor's *modularity* thesis for language.¹⁷ In the case of *formal* systems, Θ is often context free and compositional, but not always (*reader macros* can render it opaque, or at least intensional, and some languages such as ALGOL are apparently context-sensitive). It is noteworthy, however, that there have been computational languages for which Θ could not be specified independently of Ψ — a fact that is often stated as the fact that the programming language "cannot be parsed except at runtime" (TECO and the first versions of SMALLTALK had this character).
- b. Since LISP is computational, it follows that a *full* account of its Ψ can be specified independent of Φ ; this is in essence the formality condition. It is important to bring out, however, that a *local* version of Ψ will typically not be compositional in a modern computational formalism, even though such locality holds in purely extensional context-free side-effect free languages such as the λ -calculus.
- c. It is widely agreed that Ψ does not uniquely determine Φ (this is the "psychology narrowly construed" and the concomitant methodological solipsism of Putnam and Fodor and others¹⁸). However this fact is compatible with our foundational claim that computational systems are distinguished in virtue of having *some* version of Φ as part of their characterisation. A very similar point can be made for logic: although any *given* logic can (presumably) be given a mathematically-specified model theory, that theory doesn't typically tie down what is often called the *standard* model or interpretation — the interpretation that *we* use. This fact does not release us, however, from positing as a candidate logic only a formalism that humans can interpret.
- d. The declarative interpretation Φ cannot be wholly determined independent of Ψ , except in purely declarative languages (such as the λ -calculus and logic and so forth). This is to say that without some account of the effect on the processor of one fragment of a whole linguistic structure, it may be impossible to say what that processor will take another fragment as designating. The use of SETQ in LISP is an example; natural language instances will be explored, below.

This last point needs a word of explanation. It is of course possible to specify Φ in mathematical terms without any explicit mention of a Ψ -like function; the approach we use in LISP defines both Ψ and Φ in terms of the overarching function Σ mentioned above, and we could of course simply define Φ without defining Ψ at all. Our point, rather, is that any successful definition of Φ will effectively

have to do the work of Ψ , more or less explicitly, either by defining some identifiable relationship, or else by embedding that relationship within the meta-theoretic machinery. We are arguing, in other words, only that the *subject* we intend Ψ to cover must be treated in some fashion or other.

What is perhaps surprising about all of this machinery is that it must be brought to bear on a purely procedural language — *all three relationships* (Θ , Φ , and Ψ) figure crucially in an account even of LISP. We are not suggesting that LISP is *like* natural languages: to point out just one crucial difference, there is no way in LISP or in any other programming language (except PROLOG) to *say* anything, whereas the ability to say things is clearly a foundational aspect of any human language. The problem in the procedural languages is one of what we may call *assertional force*; although it is possible to construct a sentence-like expression with a clear declarative semantics (such as some equivalent of " $x = 3$ "), one cannot use it in such a way as to actually mean it — so as to have it carry any assertional weight. For example, it is trivial to set some variable x to 3, or to ask whether x is 3, but there is no way to state *that* x is 3. It should be admitted, however, that computational languages bearing assertional force are under considerable current investigation. This general interest is probably one of the reasons for PROLOG's emergent popularity; other computational systems with an explicit declarative character include for example specification languages, data base models, constraint languages, and knowledge representation languages in A.I. We can only assume that the appropriate semantics for all of these formalisms will align even more closely with an illuminating semantics for natural language.

What does all of this have to do with natural language, and with computational linguistics? The essential point is this: *if* this characterisation of formal systems is tenable, and if the techniques of standard programming language semantics can be fit into this mould, then it may be possible to combine those approaches with the techniques of programming language semantics and of logic and model theories, to construct complex and interacting accounts of Ψ and of Φ . To take just one example, the techniques that are used to construct mathematical accounts of environments and continuations might be brought to bear on the issue of dealing with the complex circumstances involving discourse models, theories of focus in dealing with anaphora, and so on; both cases involve an attempt to construct a recursively specifiable account of non-local interactions among disparate fragments of a composite text. But the contributions can proceed in the other direction as well: even from a very simple application of this framework to this circumstance of LISP, for example, we have been able to show how an accepted computational notion fails to cohere with our attributed linguistically based understanding, involving us in a major reconstruction of LISP's foundations. The similarities are striking.

Our claim, in sum, is that similar phenomena occur in programming languages and natural languages, and that each discipline could benefit from the semantical techniques developed in the other. Some examples of these similar phenomena will help to motivate this view. The first is the issue of the appropriate use of noun phrases: as well as employing a noun phrase in a standard extensional position, natural language semantics has concerned itself with more difficult cases such as *intensional contexts* (as in the underlined designator in *I didn't know The Big Apple was an island*, where the co-designating term *New York* cannot be substituted without changing the meaning), the so-called *attributive/referential*

distinction of Donellan¹⁹ (the difference, roughly, between using a noun phrase like "the man with a martini" to inform you that someone is drinking a martini, as opposed to a situation where one uses the hearer's belief or assumption that someone is drinking a martini to refer to him), and so on. Another example different from either of these is provided by the underlined term in *For the next 20 years let's restrict the president's salary to \$20,000*, on the reading in which after Reagan is defeated he is allowed to earn as much as he pleases, but his successor comes under our constraint. The analogous computational cases include for example the use of an expression like (the formal analog of) *make the sixth array element be 10* (i.e., $A(6) ::= 10$), where we mean not that the current sixth element should be 10 (the current sixth array element might at the moment be 9, and 9 can't be 10), but rather that we would like the description "the sixth array element" to refer to 10 (so-called "L-values", analogous to MACLISP's SETF construct). Or, to take a different case, suppose we say *set X to the sixth array element* (i.e., $X ::= A(6)$), where we mean not that X should be set to the current sixth array element, but that it should *always* be equal to that element (stated computationally this might be phrased as saying that X should track A(6); stated linguistically we might say that X should *mean* "the sixth array element"). Although this is not a standard type of assignment, the new constraint languages provide exactly such facilities, and macros (classic computational intensional operators) can be used in more traditional languages for such purposes. Or, for a final example, consider the standard declaration: INTEGER X, in which the term "X" refers neither to the variable *itself* (variables are *variables*, not numbers), nor to its current designation, but rather to whatever will satisfy the description "the value of X" at any point in the course of a computation. All in all, we cannot ignore the attempt on the computationalists' part to provide complex mechanisms so strikingly similar to the complex ways we use noun phrases in English.

A very different sort of linguistic phenomenon that occurs in both programming languages and in natural language are what we might call "premature exits": cases where the processing of a local fragment *aborts* the standard interpretation of an encompassing discourse. If for example I say to you *I was walking down the street that leads to the house that Mary's aunt used to ... forget it; I was taking a walk*, then the "forget it" must be used to discard the analysis of some amount of the previous sentence. The grammatical structure of the subsequent phrase determines how much has been discarded, of course; the sentence would still be comprehensible if the phrase "an old house I like" followed the "forget it". We are not accustomed to semantical theories that deal with phenomena like this, of course, but it is clear that any serious attempt to model real language understanding will have to face them. Our present point is merely that continuations²⁰ enable computational formalisms to deal exactly with the computational analogs of this: so-called *escape operators* like MACLISP's THROW and CATCH and QUIT.

In addition, a full semantics of language will want to deal with such sentences as *If by "flustrated" you mean what I think, then she was certainly flustrated*. The proper treatment of the first clause in this sentence will presumably involve lots of "ψ" sorts of considerations: its contribution to the remainder of the sentence has more to do with the mental states of speaker and hearer than with the world being describe by the presumed conversation. Once again, the overarching computational hypothesis suggests that the way these psychological effects must be modelled is in terms of alterations in

the state of an internal process running over a field of computational structures.

As well as these specific examples, a couple of more general morals can be drawn, important in that they speak directly to styles of practice that we see in the literature. The first concerns the suggestion, apparently of some currency, that we reject the notion of logical form, and "do semantics directly" in a computational model. On our account this is a mistake, pure and simple: to buy into the computational framework is to believe that the ingredients in any computational process are inherently linguistic, in need of interpretation. Thus they too will need semantics; the internalisation of English into a computer (Θ) is a translation relationship (in the sense of preserving Φ, presumably) — even if it is wildly contextual, and even if the internal language is very different in structure from the structure of English. It has sometimes been informally suggested, in an analogous vein, that Montague semantics cannot be taken seriously computationally, because the models that Montague proposes are "too big" — how could you possibly carry these infinite functions around in your head, we are asked to wonder. But of course this argument comits a use/mention mistake: the only valid computational reading of Montague would mean that mentalse (S) would consist of *designators* of the functions Montague proposes, and those designators can of course be a few short formulae.

It is another consequence of our view that any semanticist who proposes some kind of "mental structure" in his or her account of language is committed to providing an interpretation of that structure. Consider for example a proposal that posits a notion of "focus" for a discourse fragment. Such a focus might be viewed as a (possibly abstract) entity in the world, or as a element of computational structure playing such-and-such role in the behavioural model of language understanding. It might seem that these are alternative accounts: what our view insists is that an interpretation of the latter must give it a designation (Φ); thus there would be a computational structure (being biased, we will call it the *focus-designator*), and a designation (that we call the *focus-itself*). The complete account of focus would have to specify both of these (either directly, or else by relying on the generic declarative semantics to mediate between them), and also tell a story about how the focus-designator plays a causal role (ψ) in engendering the proper behaviour in the computational model of language understanding.

There is one final problem to be considered: what it is to design an internal formalism S (the task, we may presume, of anyone designing a knowledge representation language). Since, on our view, we must have a semantics, we have the option either of having the semantics informally described (or, even worse, tacitly assumed), or else we can present an explicit account, either by defining such a story ourselves or by borrowing from someone else. If the LISP case can be taken as suggestive, a purely declarative model theory will be inadequate to handle the sorts of computational interactions that programming languages have required (and there is no *a priori* reason to assume that successful computational models for natural language will be found that are *simpler* than the programming languages the community has found necessary for the modest sorts of tasks computers are presently able to perform). However it is also reasonable to expect that no direct analog to programming language semantics will suffice, since they have to date been so concerned with purely procedural (behavioural) consequence. It seems at least

reasonable to suppose that a general interpretation function, of the sort mentioned earlier, may be required.

Consider for example the KLONE language presented by Brachman *et al.*²¹ Although no semantics for KLONE has been presented, either procedural or declarative, its proponents have worked both in investigating the Θ -semantics (how to translate English into KLONE), and in developing an informal account of the procedural aspects. Curiously, recent directions in that project would suggest that its authors expect to be able to provide a "declarative-only" account of KLONE semantics (i.e., expect to be able to present an account of Φ independent of Ψ), in spite of our foregoing remarks. Our only comment is to remark that *independence* of procedural consequence is not a pre-requisite to an adequate semantics; the two can be recursively specifiable together; thus this apparent position is stronger than formally necessary — which makes it perhaps of considerable interest.

In sum, we claim that any semantical account of *either* natural language *or* computational language must specify Θ , Ψ , and Φ ; if any are left out, the account is not complete. We deny, furthermore, that there is any fundamental distinction to be drawn between so-called procedural languages (of which LISP is the paradigmatic example in A.I.) and other more declarative languages (encodings of logic, or representation languages). We deny as well, contrary to at least some popular belief, the view that a mathematically well-specified semantics for a candidate "mentalese" must be satisfied by giving an independently specified *declarative* semantics (as would be possible for an encoding of logic, for example). The designers of KRL,²² for example, for principled reasons denied the possibility of giving a semantics independent of the procedures in which the KRL structures participated; our simple account of LISP has at least suggested that such an approach could be pursued on a mathematically sound footing. Note however, in spite of our endorsement of what might be called a *procedural semantics*, that this in no way frees one from from giving a declarative semantics as well; *procedural semantics* and *declarative semantics* are two pieces of a total story; they are not alternatives.

NOTES

* I am grateful to Barbara Grosz and Hector Levesque for their comments on an earlier draft of this short paper, and to Jane Robinson for her original suggestion that it be written.

1. Smith (1982b)
2. Fodor (1978), Fodor (1980), Haugeland (forthcoming)
3. At least until the day arrives — if ever — when a successful psychology of language is presented wherein *all* of human semanticity is explained in non-semantical terms.
4. Problematic because it defines computation in a manner that is derivative on mind (in that language is fundamentally a mental phenomenon), thus dashing the hope that computational psychology will offer a release from the semantic irreducibility of previous accounts of human cognition. Though we state this position and explore some of its consequences in Smith (1982b), a considerably fuller treatment will be provided in Smith (forthcoming).
5. See for example Newell (1980)

6. The term "semantics" is only one of a large collection of terms, unfortunately, that are technical terms in computer science and in the attendant cognitive disciplines (including logic, philosophy of language, linguistics, and psychology), with different meanings and different connotations. *Reference*, *interpretation*, *memory*, and *value* are just a few examples of the others. It is our view that in spite of the fact that semantical vocabulary is used in *different* ways, the fields are both semantical in fundamentally the *same* ways: a unification of terminology would only be for the best.
7. An example of the phenomenon noted in footnote 6.
8. Fodor (forthcoming)
9. Woods (1981)
10. For a discussion of continuations see Gordon (1979), Steele and Sussman (1978), and Smith (1982a); the formal device is developed in Strachey & Wadsworth (1974).
11. Smith (1982a).
12. A classic example is Katz and Postal (1964), but much of the recent A.I. research in natural language in A.I. can be viewed in this light.
13. Lewis (1972).
14. Israel (1980).
15. For a discussion of PROLOG see Clocksin & Mellish (1981); the LISPs are described in Smith (1982a).
16. Smith (forthcoming).
17. Fodor (forthcoming).
18. The term "methodological solipsism" is from Putnam (1975); see also Fodor (1980).
19. Donnellan (1966).
20. See note 10, above.
21. Brachman (1979).
22. Bobrow and Winograd (1977).

REFERENCES

- Bobrow, Daniel G., and Winograd, Terry, "An Overview of KRL: A Knowledge Representation Language", *Cognitive Science* 1 pp. 3-46, 1977.
- Brachman, Ronald, "On the Epistemological Status of Semantic Networks", in Findler, Nicholas V. (ed.), *Associative Networks: Representation and Use of Knowledge by Computers*, New York: Academic Press, 1979.
- Clocksin, W. F., and Mellish, C. S., *Programming in Prolog*, Berlin: Springer-Verlag, 1981.
- Donnellan, K., "Reference and Definite Descriptions", *Philosophical Review* 75:3 (1966) pp. 281-304; reprinted in Rosenberg and Travis (eds.), *Readings in the Philosophy of Language*, Prentice-Hall, 1971.
- Fodor, Jerry, "Tom Swift and his Procedural Grandmother", *Cognition* 6, 1978; reprinted in Fodor (1981).
- , "Methodological Solipsism Considered as a Research Strategy in Cognitive Psychology", *The Behavioral and Brain Sciences*, 3:1 (1980) pp. 63-73; reprinted in Haugeland (ed.), *Mind Design*, Cambridge: Bradford, 1981, and in Fodor (1981).

- Israel, David, "What's Wrong with Non-Monotonic Logic?", Proceedings of the First Annual Conference of the American Association for Artificial Intelligence, Stanford, California, 1980, pp. 99-101.
- Katz, Jerrold, and Postal, Paul, *An Integrated Theory of Linguistic Descriptions*, Cambridge: M.I.T. Press, 1964.
- Lewis, David, "General Semantics", in Davidson and Harman (eds.), *Semantics of Natural Languages*, Dordrecht, Holland: D. Reidel, 1972, pp. 169-218.
- Newell, Allen, "Physical Symbol Systems", *Cognitive Science* 4, pp. 135-183, 1980.
- Putnam, Hilary, "The meaning of 'meaning'", in Putnam, Hilary, *Mind, Language and Reality*, Cambridge, U.K.: Cambridge University Press, 1975.
- Smith, Brian C., *Reflection and Semantics in a Procedural Language*, Laboratory for Computer Science Report LCS-TR-272, M.I.T., Cambridge, Mass., 1982 (a).
- , "Semantic Attribution and the Formality Condition", presented at the Eighth Annual Meeting of the Society for Philosophy and Psychology, London, Ontario, Canada, May 13-16, 1982 (b).
- , *The Computational Metaphor*, Cambridge: Bradford (forthcoming).
- Steele, Guy, and Sussman, Gerald J., "The Art of the Interpreter, or the Modularity Complex (parts Zero, One, and Two)", M.I.T. Artificial Intelligence Laboratory Memo AIM-453, Cambridge, Mass, 1978.
- Strachey, C., and Wadsworth, C. P., "Continuations — a Mathematical Semantics for Handling Full Jumps", PRG-11, Programming Research Group, University of Oxford, 1974.
- Woods, William A., "Procedural Semantics as a Theory of Meaning", Report No. 4627, Bolt Beranek and Newman, 50 Moulton St., Cambridge, Mass., 02138; reprinted in Joshi, A., Sag, I., and Webber, B., *Computational Aspects of Linguistic Structures and Discourse Settings*, Cambridge, U.K.: Cambridge University Press, 1982.